



Motivation



In the age of "Filter Bubble", readers who have not heard of the controversy is likely to be misled or uninformed.



Watch the full Youtube Clip "If Google Was A Guy"

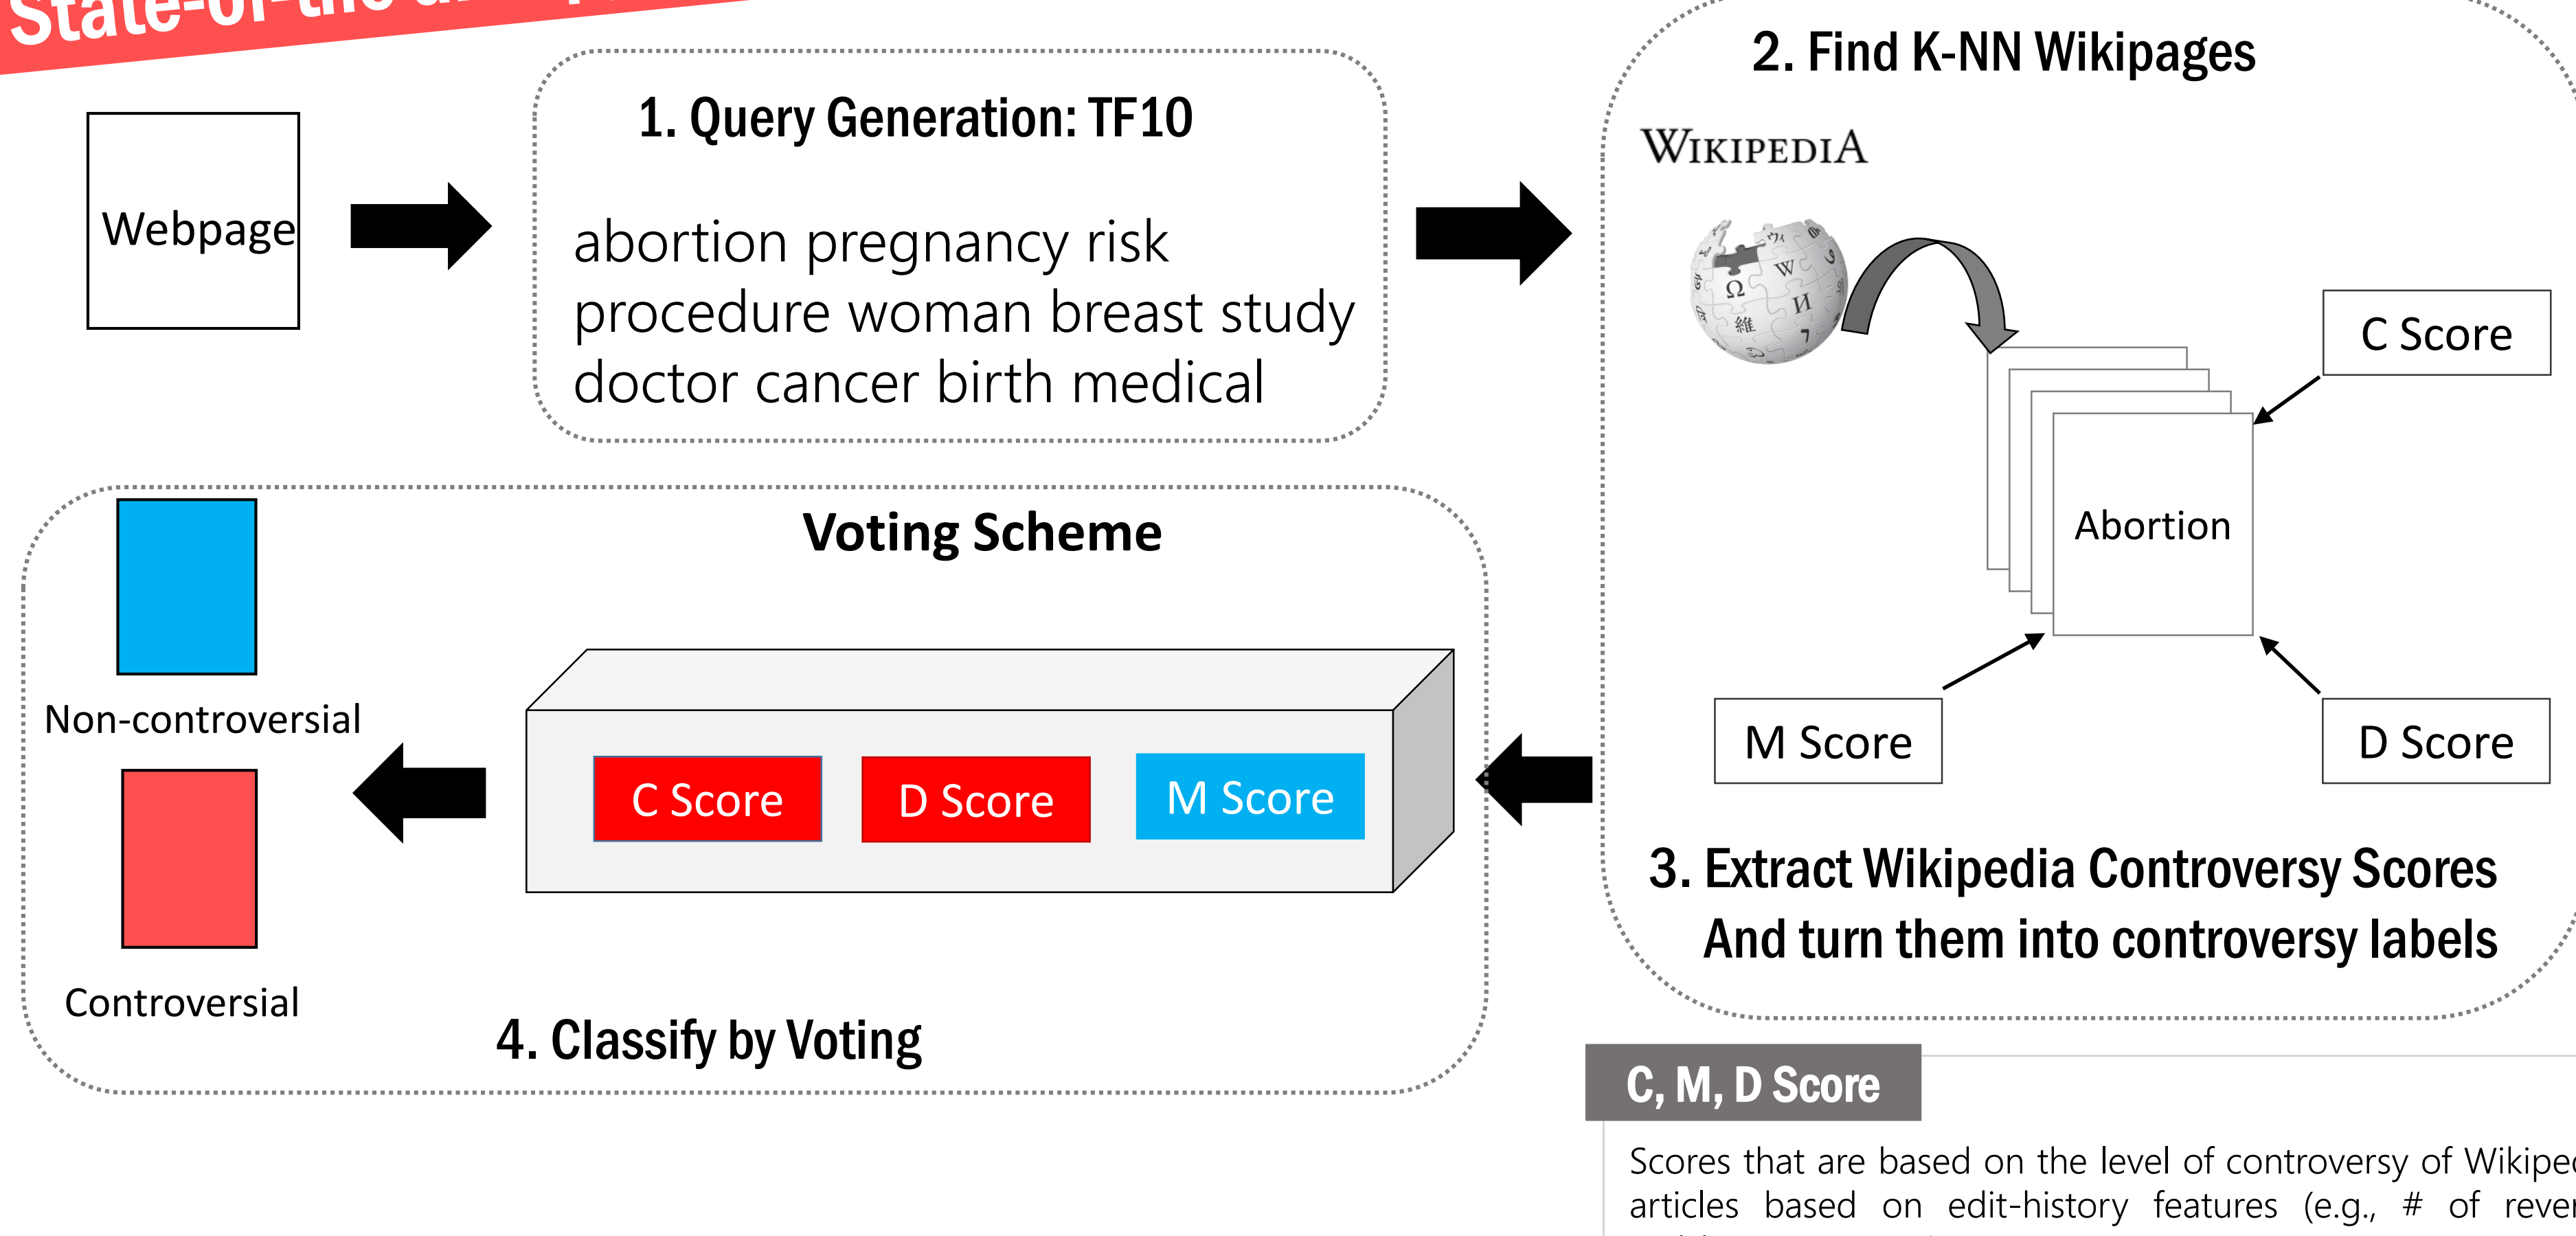
Research Goal

Automatically warning users when a Web-page discusses controversial topics



Controversial

State-of-the art Approach



1. Query Generation: TF10
abortion pregnancy risk procedure woman breast study doctor cancer birth medical

2. Find K-NN Wikipages

3. Extract Wikipedia Controversy Scores And turn them into controversy labels

4. Classify by Voting

C Score, D Score, M Score

Non-controversial, Controversial


S. Dori-Hacohen and J. Allan, Automated Controversy Detection on the Web, ECIR 2015

Limitations

- Ambiguous Query Generation**
The generated query ends up containing an unknown mixture of different sub-topics of a document.
➔ **Solution 1: Generating multiple queries from smaller but more coherent blocks (or tiles) of the document**
- Underrated Wikipedia Controversy Scores**
The controversy level is underestimated on specific and sub-topical Wikipages whose topical disputes have often been delegated by other Wikipages on the broader topic.
➔ **Solution 2. Smoothing Wikipedia controversy scores using neighbors**

Solution 1: Finding KNN Pages with TileQuery

- Document Segmentation**
Split a document into multiple block of sentences (tiles) using *TextTiling*.
- Query Generation**
Generate a query of a mixture of g (global) and l (local) most frequent terms from the tile.
- Aggregating the Ranked Lists**
Compute the relevance of each retrieved Wikipedia page w_i by aggregating its rank scores as follows:



Experiments

Dataset

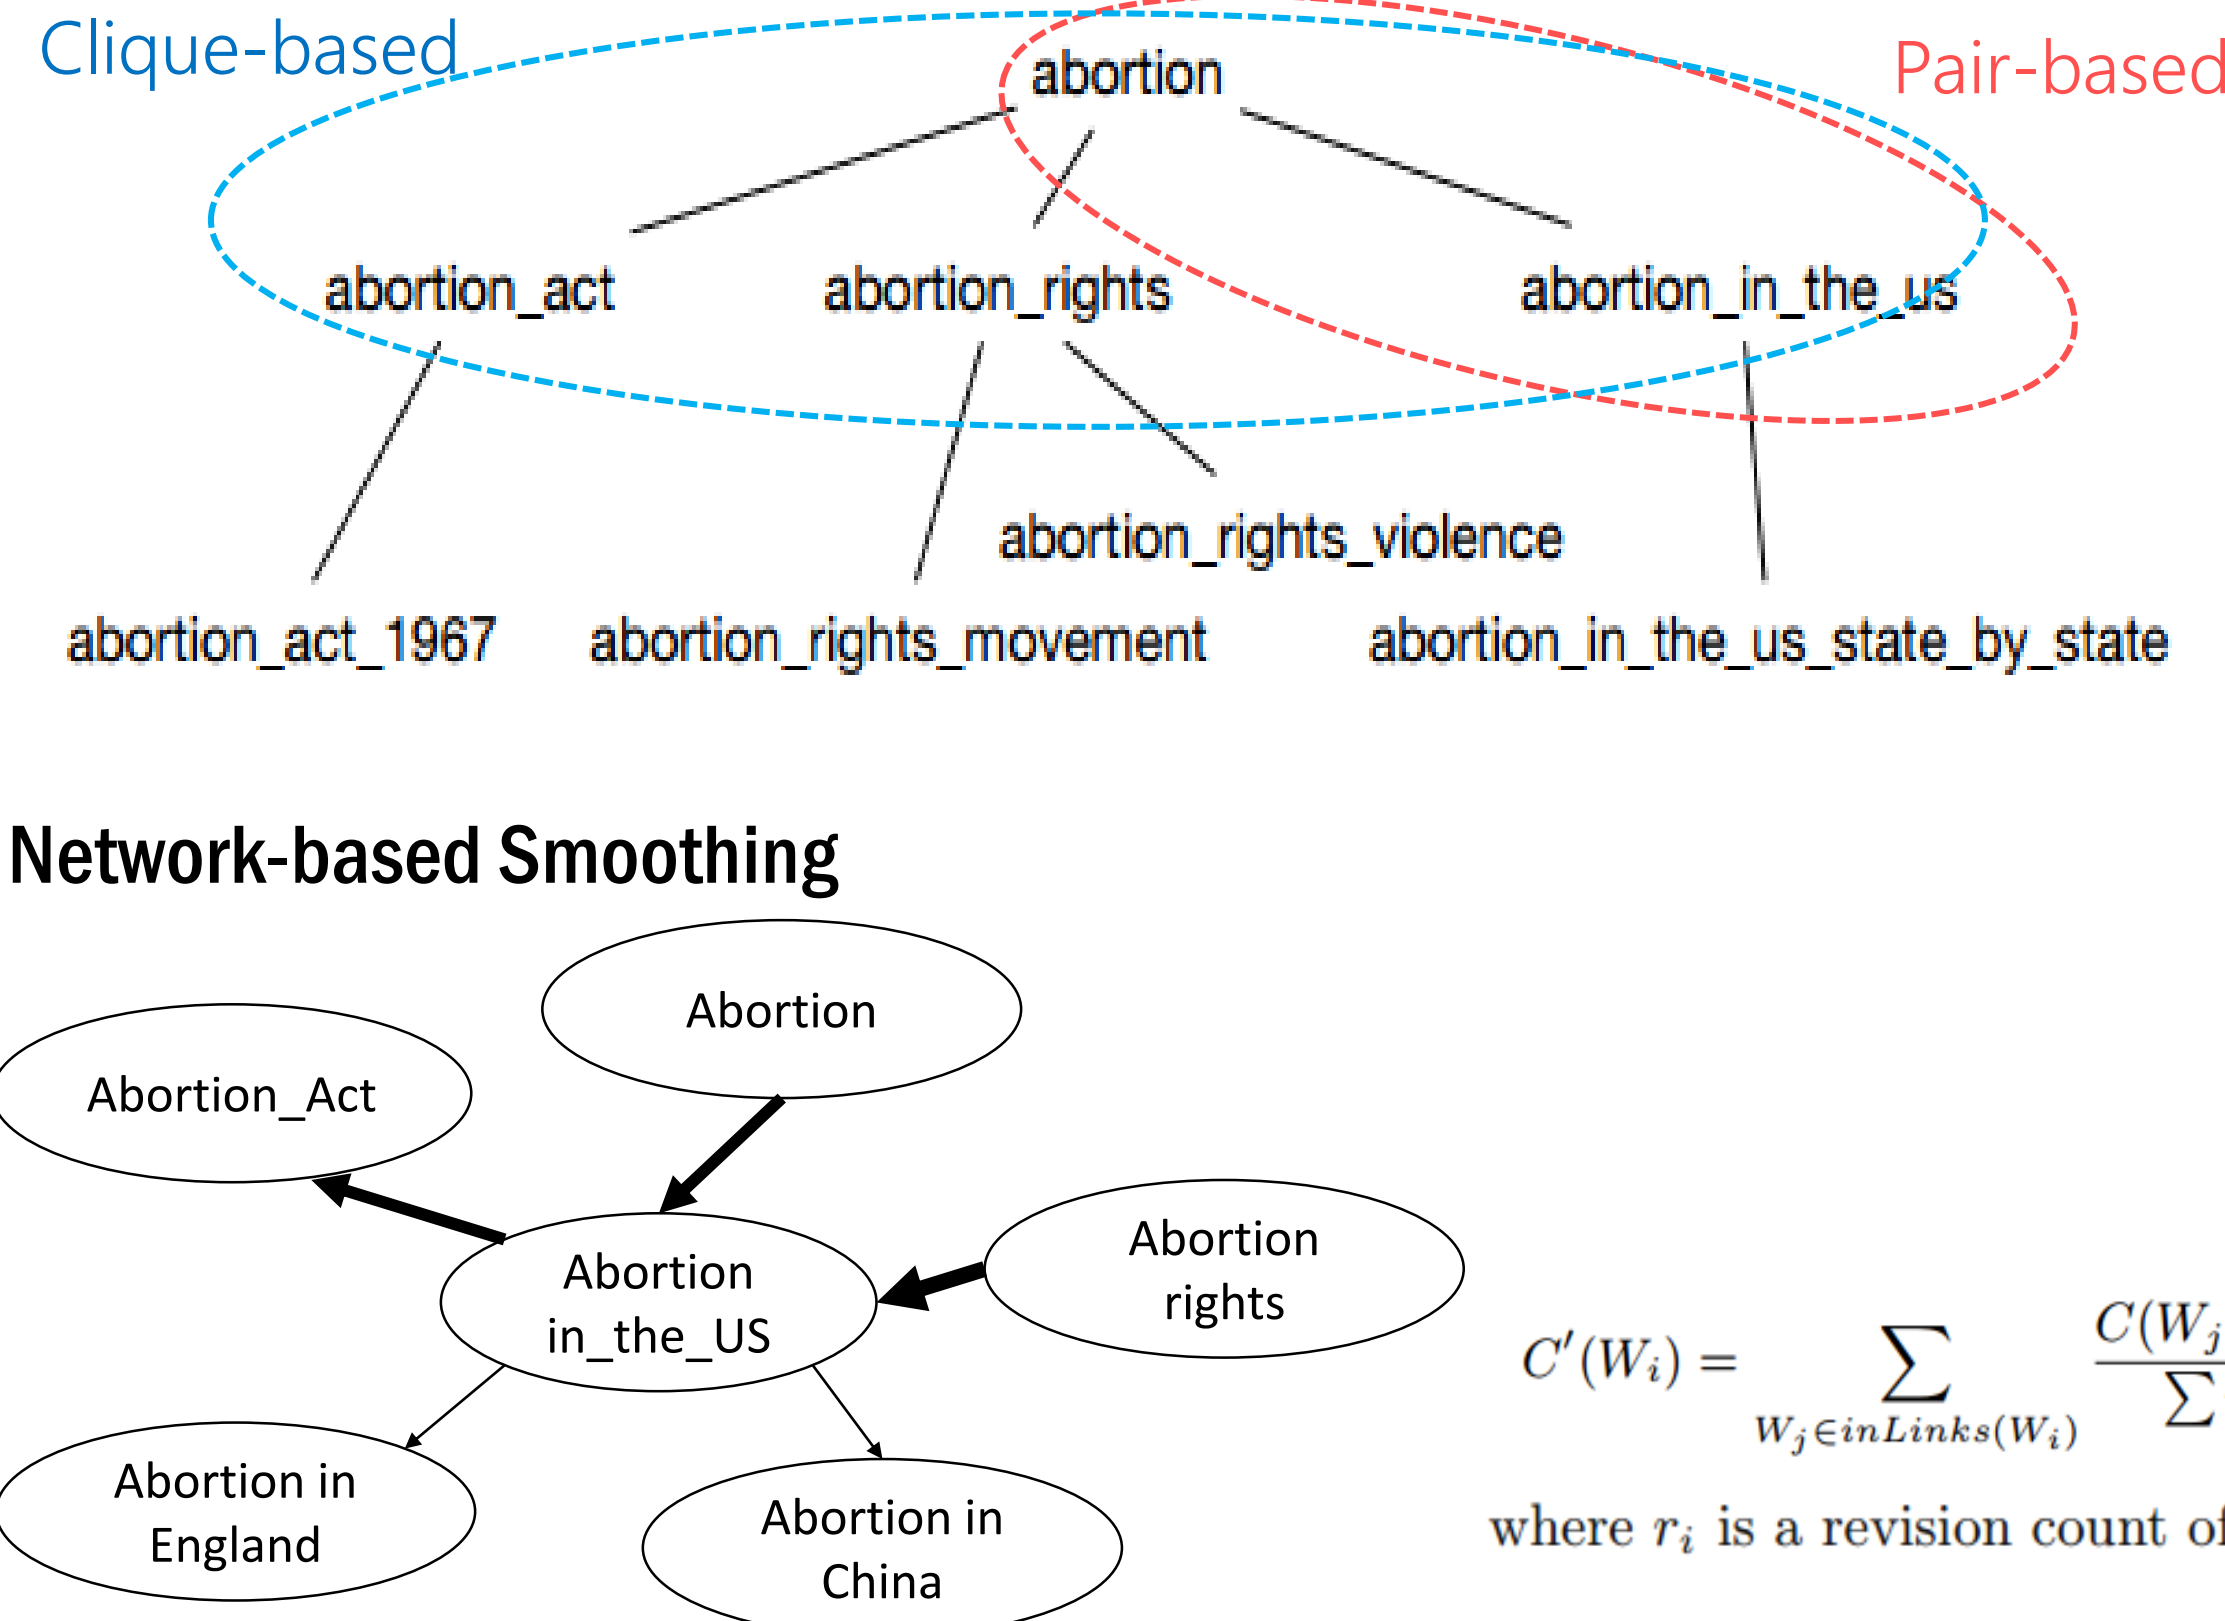
- 303 Web documents
- Collected from topics with varying controversy level
- 42%: controversial

Extensive Parameter Explorations:

- Query Methods = {AllQuery, TF10, TileQuery}
- # of Wikipedia Neighbors, $K = \{1... 20\}$
- Use of score smoothing
- Thresholds for C and M Scores
- Wikipedia neighbors selection = {Pair, Clique}
- Voting Method = {M, C, D, Majority, Or/AND}

Solution 2: Smoothing Wikipedia Controversy Scores

- Construct Wikipedia Topic Hierarchy using their titles' prefix-relation**
Clique-based, Pair-based
- Network-based Smoothing**



$$C'(W_i) = \sum_{W_j \in \text{inLinks}(W_i)} \frac{C(W_j) * r_j}{\sum r_k}$$

where r_i is a revision count of W_i .

	Original scores		Revised scores	
	M	C	M	C
Abortion	4,102,593	0.300	4,102,593	0.300
Abortion_Act	0	0	0	0
Abortion_in_China	0	0	2,062,156	0.166
Abortion_in_England	0	0	2,128,909	0.172
Abortion_in_the_US	0	0.002	2,983,300	0.218

Two controversy scores on "abortion"-related pages, before and after smoothing the score

Table 2: Accuracy, F1, and the best parameters in 5-fold runs for different query and inferred score settings.

Run	Query	Inferred Score	K	C Threshold	M Threshold	Aggregation	Acc.	F1
1		N/A	{5, 20}	{0.17, 4.18 · 10 ⁻² }	{40000,20000}	{M, Maj.}	0.72	0.50
2	ALL	Clique	15	{0.17, 4.18 · 10 ⁻² }	{40000,20000}	{M, Maj.}	0.78	0.68
3		Pair	{5, 20}	{0.17, 4.18 · 10 ⁻² }	{40000,20000}	{M, Maj.}	0.73	0.53
4		N/A	20	4.18 · 10 ⁻²	{20000, 40000, 84930}	{M, Maj.}	0.75	0.57
5	TF10	Clique	20	4.18 · 10 ⁻²	84930	Maj.	0.79	0.68
6		Pair	{10, 20}	4.18 · 10 ⁻²	{20000, 84930}	Maj.	0.75	0.57
7		N/A	{10,15,20}	4.18 · 10 ⁻²	{40000,20000}	{M, Maj.}	0.75	0.59
8	TILE	Clique	20	0.17	40000	M	0.80	0.71
9		Pair	{10,15,20}	4.18 · 10 ⁻²	{40000,20000}	{M, Maj.}	0.75	0.61

Conclusion and Future Work

- Our modifications of TileQuery and Wikipedia controversy score smoothing improved the state-of-the-art controversy detection by 5% in Acc and 14% in F1.
- Title prefix-relation only covered 5~10% of the Wikipedia titles. More sophisticated methods to find Wikipedia neighbors for smoothing can be studied.
- We plan to further investigate different scenarios where TF10 and TileQuery works well.